

# Practical 1: Speeding up performance with openmp

GAHAN SARAIYA (18MCEC10)

[18mcec10@nirmauni.ac.in](mailto:18mcec10@nirmauni.ac.in)

## I. AIM

- Speeding up performance with openmp

## II. INTRODUCTION

Generic steps:

- Making code sequential code to parallel using openmp
- Do speedup calculation – Sequential/parallel
- compute the same speedup for various values

## III. IMPLEMENTATION

**Implementing Vector Addition** Steps:

- Implement Sequential code
- Implement Parallel code
- Implement function for verification of output (for accuracy calculation)
- Observe time for consumed for execution of sequential and parallel code over 100 iterations (repeating operation for 100 times)
- Calculate speedup for various input size

### I. Sequential Code

```
// -----  
// Author: Gahan Saraiya  
// GiT: http://github.com/gahan9/  
// StackOverflow: https://stackoverflow.com/users/story/7664524  
// Website: http://gahan9.github.io/  
// -----  
// Making code sequential code to parallel using openmp  
// Do speedup calculation -- Sequential/parallel  
// Show profiling and total execution time  
// Standard vector addition (serial, without OMP)
```

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <omp.h>

#define ARRAY_SIZE 1000000000
#define REPEAT    100

void vector_add(double* vector1, double* vector2, double* result) {
    for (int i=0; i < ARRAY_SIZE; i++){
        result[i] = vector1[i] + vector2[i];
    }
}

double* generate_array(int no_of_elements) {
    // no_of_elements : Generate array of `no_of_elements` number of elements
    double* array = (double*) malloc(no_of_elements*sizeof(double));
    for(int i=0; i<no_of_elements; i++){
        array[i] = rand()%10000;
    }
    return array;
}

int verify(double* vector1, double* vector2) {
    // Verifying accuracy of computation
    double *adder = (double*) malloc(ARRAY_SIZE*sizeof(double));
    double *verifier = (double*) malloc(ARRAY_SIZE*sizeof(double));
    vector_add(vector1, vector2, adder);

    for(int i=0; i<ARRAY_SIZE; i++) {
        verifier[i] = vector1[i] + vector2[i];
    }
    for(int i=0; i<ARRAY_SIZE; i++){
        if(verifier[i] != adder[i]){
            return 0;
        }
    }
    return 1;
}

int main() {
    // Generate input vectors and destination vector
    double *vector1 = generate_array(ARRAY_SIZE);
    double *vector2 = generate_array(ARRAY_SIZE);
    double *result_vector = (double*) malloc(ARRAY_SIZE*sizeof(double));
```

```
// Double check vector_add is correct
if(!verify(vector1, vector2)) {
    printf("vector_add does not match oracle\n");
    return 0;
}

// Test framework that sweeps the number of threads and times each
// runs for iteration REPEAT
double start_time, run_time;
int num_threads = 1;
for(int i=1; i<=num_threads; i++) {
    omp_set_num_threads(i);
    start_time = omp_get_wtime();
    for(int j=0; j<REPEAT; j++){
        vector_add(vector1, vector2, result_vector);
    }
    run_time = omp_get_wtime() - start_time;
    printf(" %d thread(s) took %f seconds\n",i,run_time);
}
}
```

## II. Parallel Code

```
// -----
// Author: Gahan Saraiya
// GiT: http://github.com/gahan9/
// StackOverflow: https://stackoverflow.com/users/story/7664524
// Website: http://gahan9.github.io/
// -----
// Making code sequential code to parallel using openmp
// Do speedup calculation -- Sequential/parallel
// Show profiling and total execution time
// Standard vector addition (serial, without OMP)

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <omp.h>

#define ARRAY_SIZE 100000000
#define REPEAT 100
#define NUM_THREADS 4

void vector_add(double* vector1, double* vector2, double* result) {
    int i;
    int size_chunks = ARRAY_SIZE/NUM_THREADS;
    #pragma omp parallel
    {
```

```
// #pragma omp parallel for num_threads(NUM_THREADS)
// #pragma omp parallel for
// for (i=0; i < ARRAY_SIZE; i++){
    for(int i=size_chunks*omp_get_thread_num(); i < size_chunks*(omp_get_thread_num()+1); i++)
        result[i] = vector1[i] + vector2[i];
}
}

double* generate_array(int no_of_elements) {
    // no_of_elements : Generate array of `no_of_elements` number of elements
    double* array = (double*) malloc(no_of_elements*sizeof(double));
    for(int i=0; i<no_of_elements; i++)
        array[i] = rand()%10000;
    return array;
}

int verify(double* vector1, double* vector2) {
    // Verifying accuracy of computation
    double *adder = (double*) malloc(ARRAY_SIZE*sizeof(double));
    double *verifier = (double*) malloc(ARRAY_SIZE*sizeof(double));
    vector_add(vector1, vector2, adder);

    for(int i=0; i<ARRAY_SIZE; i++) {
        verifier[i] = vector1[i] + vector2[i];
    }
    for(int i=0; i<ARRAY_SIZE; i++){
        if(verifier[i] != adder[i]){
            return 0;
        }
    }
    return 1;
}

int main() {
    // Generate input vectors and destination vector
    double *vector1 = generate_array(ARRAY_SIZE);
    double *vector2 = generate_array(ARRAY_SIZE);
    double *result_vector = (double*) malloc(ARRAY_SIZE*sizeof(double));

    // Double check vector_add is correct
    if(!verify(vector1, vector2)) {
        printf("vector_add does not match actual result\n");
        return 0;
    }
}
```

```

    // Test framework that sweeps the number of threads and times each
    // runs for iteration REPEAT
    double start_time, run_time;
    start_time = omp_get_wtime();

    for(int j=0; j<REPEAT; j++){
        vector_add(vector1, vector2, result_vector);
    }
    run_time = omp_get_wtime() - start_time;
    printf(" %d thread(s) took %f seconds\n", NUM_THREADS, run_time);
}

```

#### IV. ANALYSIS

Below is the result of observation over 100 iterations for various input size

Data size	Sequential Time	Parallel Time <sub>1</sub>	Parallel Time <sub>2</sub>	Speedup <sub>1</sub>	Speedup <sub>2</sub>
10 <sup>5</sup>	0.003093	0.002066	0.00227	1.497095837	1.362555066
10 <sup>6</sup>	0.031473	0.029124	0.022613	1.08065513	1.391810021
10 <sup>7</sup>	0.366141	0.305616	0.313129	1.198042642	1.169297638
10 <sup>8</sup>	3.490682	3.014184	2.952571	1.15808524	1.182251672
10 <sup>9</sup>	35.395422	29.491622	29.188446	1.200185666	1.212651814




Figure 1: Analysis of Time consumed by thread(s) for calculation of number of inputs (x-axis scaled to the log base 10)




Figure 2: Comparison of speed up achieved from two method of parallelism