

Practical 1: Speeding up performance with openmp

GAHAN SARAIYA (18MCEC10)

18mcec10@nirmauni.ac.in

I. AIM

- Speeding up performance with openmp

II. INTRODUCTION

Generic steps:

- Making code sequential code to parallel using openmp
- Do speedup calculation – Sequential/parallel
- compute the same speedup for various values

III. IMPLEMENTATION

Implementing Vector Addition Steps:

- Implement Sequential code
- Implement Parallel code
- Implement function for verification of output (for accuracy calculation)
- Observe time for consumed for execution of sequential and parallel code over 100 iterations (repeating operation for 100 times)
- Calculate speedup for various input size

I. Sequential Code

```
// -----  
// Author: Gahan Saraiya  
// GiT: http://github.com/gahan9/  
// StackOverflow: https://stackoverflow.com/users/story/7664524  
// Website: http://gahan9.github.io/  
// -----  
// Making code sequential code to parallel using openmp  
// Do speedup calculation -- Sequential/parallel  
// Show profiling and total execution time  
// Standard vector addition (serial, without OMP)
```

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <omp.h>

#define ARRAY_SIZE 1000000000
#define REPEAT    100

void vector_add(double* vector1, double* vector2, double* result) {
    for (int i=0; i < ARRAY_SIZE; i++){
        result[i] = vector1[i] + vector2[i];
    }
}

double* generate_array(int no_of_elements) {
    // no_of_elements : Generate array of `no_of_elements` number of elements
    double* array = (double*) malloc(no_of_elements*sizeof(double));
    for(int i=0; i<no_of_elements; i++){
        array[i] = rand()%10000;
    }
    return array;
}

int verify(double* vector1, double* vector2) {
    // Verifying accuracy of computation
    double *adder = (double*) malloc(ARRAY_SIZE*sizeof(double));
    double *verifier = (double*) malloc(ARRAY_SIZE*sizeof(double));
    vector_add(vector1, vector2, adder);

    for(int i=0; i<ARRAY_SIZE; i++) {
        verifier[i] = vector1[i] + vector2[i];
    }
    for(int i=0; i<ARRAY_SIZE; i++){
        if(verifier[i] != adder[i]){
            return 0;
        }
    }
    return 1;
}

int main() {
    // Generate input vectors and destination vector
    double *vector1 = generate_array(ARRAY_SIZE);
    double *vector2 = generate_array(ARRAY_SIZE);
    double *result_vector = (double*) malloc(ARRAY_SIZE*sizeof(double));
```

```
// Double check vector_add is correct
if(!verify(vector1, vector2)) {
    printf("vector_add does not match oracle\n");
    return 0;
}

// Test framework that sweeps the number of threads and times each
// runs for iteration REPEAT
double start_time, run_time;
int num_threads = 1;
for(int i=1; i<=num_threads; i++) {
    omp_set_num_threads(i);
    start_time = omp_get_wtime();
    for(int j=0; j<REPEAT; j++){
        vector_add(vector1, vector2, result_vector);
    }

    run_time = omp_get_wtime() - start_time;
    printf(" %d thread(s) took %f seconds\n",i,run_time);
}
}
```

II. Parallel Code

```
// -----
// Author: Gahan Saraiya
// GiT: http://github.com/gahan9/
// StackOverflow: https://stackoverflow.com/users/story/7664524
// Website: http://gahan9.github.io/
// -----
// Making code sequential code to parallel using openmp
// Do speedup calculation -- Sequential/parallel
// Show profiling and total execution time
// Standard vector addition (serial, without OMP)

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <omp.h>

#define ARRAY_SIZE 100000000
#define REPEAT 100
#define NUM_THREADS 4

void vector_add(double* vector1, double* vector2, double* result) {
    int i;
    int size_chunks = ARRAY_SIZE/NUM_THREADS;
    #pragma omp parallel
    {
```

```
// #pragma omp parallel for num_threads(NUM_THREADS)
// #pragma omp parallel for
// for (i=0; i < ARRAY_SIZE; i++){
    for(int i=size_chunks*omp_get_thread_num(); i < size_chunks*(omp_get_thread_num()+1); i++)
        result[i] = vector1[i] + vector2[i];
}
}

double* generate_array(int no_of_elements) {
    // no_of_elements : Generate array of `no_of_elements` number of elements
    double* array = (double*) malloc(no_of_elements*sizeof(double));
    for(int i=0; i<no_of_elements; i++)
        array[i] = rand()%10000;
    return array;
}

int verify(double* vector1, double* vector2) {
    // Verifying accuracy of computation
    double *adder = (double*) malloc(ARRAY_SIZE*sizeof(double));
    double *verifier = (double*) malloc(ARRAY_SIZE*sizeof(double));
    vector_add(vector1, vector2, adder);

    for(int i=0; i<ARRAY_SIZE; i++) {
        verifier[i] = vector1[i] + vector2[i];
    }
    for(int i=0; i<ARRAY_SIZE; i++){
        if(verifier[i] != adder[i]){
            return 0;
        }
    }
    return 1;
}

int main() {
    // Generate input vectors and destination vector
    double *vector1 = generate_array(ARRAY_SIZE);
    double *vector2 = generate_array(ARRAY_SIZE);
    double *result_vector = (double*) malloc(ARRAY_SIZE*sizeof(double));

    // Double check vector_add is correct
    if(!verify(vector1, vector2)) {
        printf("vector_add does not match actual result\n");
        return 0;
    }
}
```

```

    // Test framework that sweeps the number of threads and times each
    // runs for iteration REPEAT
    double start_time, run_time;
    start_time = omp_get_wtime();

    for(int j=0; j<REPEAT; j++){
        vector_add(vector1, vector2, result_vector);
    }
    run_time = omp_get_wtime() - start_time;
    printf(" %d thread(s) took %f seconds\n", NUM_THREADS, run_time);
}

```

IV. ANALYSIS

Below is the result of observation over 100 iterations for various input size

Data size	Sequential Time	Parallel Time ₁	Parallel Time ₂	Speedup ₁	Speedup ₂
10 ⁵	0.003093	0.002066	0.00227	1.497095837	1.362555066
10 ⁶	0.031473	0.029124	0.022613	1.08065513	1.391810021
10 ⁷	0.366141	0.305616	0.313129	1.198042642	1.169297638
10 ⁸	3.490682	3.014184	2.952571	1.15808524	1.182251672
10 ⁹	35.395422	29.491622	29.188446	1.200185666	1.212651814




Figure 1: Analysis of Time consumed by thread(s) for calculation of number of inputs (x-axis scaled to the log base 10)




Figure 2: Comparison of speed up achieved from two method of parallelism

V. CONCLUSION

To get the speedup accurately the data and graphs values are based on summation of 100 iterations and it's been observed that for very initial where input vector size is about 10^7 then observed amount of time with parallel implementation is affecting negligible enhancement, however as the input size increase afterwards overall speedup observed about 1.21.